

Chunked

Chunked processing with dplyr

Edwin de Jonge, Statistics Netherlands
@edwindjonge | github.com/edwindj

Who am I?

- ▶ Data scientist / Methodologist at Statistics Netherlands (aka CBS).
- ▶ Author of several R-packages, including `whisker`, `validate`, `errorlocate`, `docopt`, `daff`, `tableplot`, `ffbase`,...
- ▶ Co-author of *Statistical Data Cleaning with applications in R* (2018) (together with @markvdloo)
- ▶ Co-worker of next speaker (Jan van der Laan)

What is chunked?

Short answer:



for data in text files

Process Data in GB-sized Text Files:

(pre)Process text files to:

- ▶ select columns
- ▶ filter rows
- ▶ derive new variables



Save result into:

- ▶ Another text file
- ▶ A database

Option 1: Read data with R

Use:

- ▶ `read.csv uh, readr::read_csv1`
- ▶ `datatable::fread`
- ▶ Fast reading of data into memory!

However...

- ▶ You will need a lot of RAM!
- ▶ Text files tend to be 1 to 100 Gb.
- ▶ **Even though these procedures use memory mapping the resulting data.frame does not!**
- ▶ development cycle of processing script is looooooong...

Option 2: Use unix tools

Good choice!

- ▶ sed
- ▶ awk
- ▶ grep
- ▶ fast processing!

However...

It is nice to stay in R-universe (one data-processing tool)

- ▶ Instead of learning at least 3 extra tools sed, awk and grep voodoo.
- ▶ Does it work on my OS/shell?
- ▶ I want to use dplyr verbs! (dplyr-deprivation...)

Option 3: Import data in DB

Import data into DB

- ▶ Use DB tool to import data.
- ▶ Process database with dplyr.

However

- ▶ It is not really a R, but a DB solution
- ▶ May be not efficient.

Process in chunks?



Option 4: Use chunked!

Idea:

- ▶ Process data chunk by chunk using `dplyr` verbs
 - ▶ Memory efficient, only one chunk at a time in memory
 - ▶ Lazy processing
 - ▶ Development cycle is short: test on first chunk.
-
- ▶ Read (and write) on chunk at a time using R package `LaF`.
 - ▶ All `dplyr` verbs on `chunk_wise` objects are recorded and replayed when writing.

Scenario 1: TXT -> TXT

Preprocess a text file with data

```
read_chunkwise("my_data.csv", chunk_size = 5000) %>%  
  select(col1, col2) %>%  
  filter(col1 > 1) %>%  
  mutate(col3 = col1 + 1) %>%  
write_chunkwise("output.csv")
```

This code:

- ▶ evals chunk by chunk
- ▶ allows for column name completion in Rstudio!

Scenario 2: TXT -> DB

Insert processed text data in DB

```
db <- src_sqlite('test.db', create=TRUE)

tbl <-
  read_chunkwise("./large_file_in.csv") %>%
  select(col1, col2, col5) %>%
  filter(col1 > 10) %>%
  mutate(col6 = col1 + col2) %>%
  write_chunkwise(db, 'my_large_table')
```

Scenario 3: DB -> TXT

Extract a large table from a DB to a text file

```
tbl<-  
  ( src_sqlite("test.db") %>%  
    tbl("my_table")  
  ) %>%  
  read_chunkwise(chunk_size=5000) %>%  
  select(col1, col2, col5) %>%  
  filter(col1 > 10) %>%  
  mutate(col6 = col1 + col2) %>%  
  write_chunkwise('my_large_table.csv')
```

Caveat

Working:

- ▶ Working on chunks is memory efficient
- ▶ `filter`, `select`,
`rename`, `mutate`, `mutate_each`, `transmute`, `do`, `tbl_vars`,
`inner_join`, `left_join`, `semi_join`, `anti_join` all work ,
also with name completion!

However:

- ▶ **`summarize` and `group_by` work chunkwise (and not for all data!)**
- ▶ `No arrange`, `right_join`, `full_join`

Usage?

Thank you!

Interested?

```
install.packages("chunked")
```

Ideas and suggestions?

<http://github.com/edwindj/chunked>